



How to Train Your AI Dragon: Email #6 - Job Hunting

Hello, and welcome to the last email in our short series for social scientist! (Was this forwarded? [Sign up here](#)).

In this email, I'm going to share what I have learned about job hunting with AI. I was laid off from my last job in October 2025, after being there for 16 years. After taking some time to grieve, I started researching the new job-hunting landscape.

The good news is, there are some really great tools out there now. The bad news is that there are many, many people chasing the same small pool of jobs, and it's a really difficult time to be looking.

But fear not! This series has given you a secret weapon: you know how to train your AI dragon.

"Everything We Know About Them is Wrong"

Early in the movie *How to Train Your Dragon*, Hiccup reads the Dragon Manual. That's the village's definitive guide to hunting, fighting, and surviving dragons. Every entry ends with a blunt directive: **Extreme danger. Kill on sight.**

The turning point happens when Hiccup interacts with Toothless and realizes the manual is useless. He utters the line that shifts the entire story: "Everything we know about them is wrong."

The Job Hunting Tie-In

Most traditional job-hunting advice is the equivalent of the old Dragon Manual. It tells you to send out hundreds of resumes, use keyword-stuffing hacks to beat the applicant tracking system (ATS), and treat the employer-side AI as a hostile beast that you have to outsmart or conquer.

When you treat AI or the market like an adversary to hunt down, you just end up exhausted. The modern trick is realizing that these tools aren't enemies to fight.

How AI is Used on the Employer Side

Many recruiting teams are shifting away from keyword-matching systems toward complex, multi-layered workflows and autonomous agents.

- **Semantic Screening & Contextual Matching:** Modern ATS systems like [Workable](#) use *semantic matching*. They don't just look for the word Python, they look at the surrounding context to understand *how* you used it and your level of autonomy.
- **Candidate Sourcing Agents:** Platforms like [Pin](#) use automated models to scan massive, public profile databases (spanning hundreds of millions of professional profiles) to proactively score and rank passive candidates against an open position before a recruiter even looks at an active application pile.

- **Interaction Logging:** Tools like [Read AI](#) are increasingly used to transcribe interviews, analyze technical discussions, and log context into centralized talent databases, replacing human notes with searchable transcripts.

Tailoring Resumes & Cover Letters

The biggest mistake job seekers make is asking an LLM to *"write a resume for this job description."* This yields generic, AI-sounding prose that recruiters spot instantly. Instead, build a **Fact Bank** that treats your real career history as the absolute source of truth.

Setting Up Your System

- **Establish a Fact Bank:** Create a single document containing every project, metric, tool, and role you have ever had. Include extensive, unedited details.
- **Construct a Strict System Prompt:** When passing your resume and a target job description to an AI tool, enforce strict boundaries. You can use an open-source framework or prompt an LLM directly using this logical structure:

*You are an expert resume strategist. I am going to provide a Fact Bank of my career and a target Job Description. Rewrite and reorder my resume to emphasize the experiences most relevant to the target role. **Strict Rule:** Change emphasis, phrasing, and context, but do not invent new skills, metrics, dates, or achievements. If something isn't in my Fact Bank, do not include it.*

- **Open-Source Frameworks:** If you want to see an example, look at open-source architectures like [JaimeYeung/Resume-Tailor-AI on GitHub](#). It programmatically extracts your history into a Fact Bank and uses a keyword report to adjust your existing bullet points without fabricating data. Requires an OpenAI API key.
- Another excellent local, privacy-focused tool is [rotsl/resume-tailor on GitHub](#). Requires Notion plus an Anthropic (Claude) or Google (Gemini) local API key to handle the heavy lifting.

Testing Your Resume Against the Bots

If you want to run a diagnostic to see how an automated system reads your profile, there are dedicated public tools that reverse-engineer common

recruiter ATS parsing algorithms.

- **Jobscan:** The [Jobscan ATS Resume Checker](#) allows you to paste your resume alongside a specific target job description. It generates a weighted match report identifying missing hard skills, tool proficiencies, and structural red flags (like complex multi-column layouts that break older text parsers).
- **LinkedIn:** This Microsoft tool has a [premium](#) feature that will highlight jobs for which you might be a top candidate and do tailored job searches, write a cover letter, and more. Premium costs a fair penny, but this tool is a one-stop-shop for job hunting.
- **Cake AI:** For a lighter, alternative evaluation tool, the [Cake AI Resume Checker](#) scores resume formatting, keyword density, and bullet-point impact against a specific job posting, offering direct inline suggestions for refinement.

Job Aggregators & Dashboards

Using AI to scrape listings and pipe them into a centralized hub is a great alternative to endlessly scrolling through job boards. As I've mentioned before, my own preferred tool (the only one I pay for) is [Taskade](#).

One of the first things I used Taskade for was to develop my own tailored job listing scraper, which I call my **Morning Job Scout**. It is set to scrape job listing sites every morning and put them on a dashboard. The live app is at <https://morning-job-scout-4984.taskade.app>, and the template is at <https://www.taskade.com/share/apps/7pc58z0t8fp9cifq>.

While my scout is tailored for my job search, there are a couple of public-facing templates and architectures that can help you design your own scout.

- **Notion Hubs:** The [Notion Marketplace Jobs Dashboard](#) is a widely used public layout. It provides structured relational databases to map companies, roles, URLs, and application timelines.
- **The AI Scraping Pipeline:** To make your own scraper, you'll need a web-scraping script (or a zero-code scraper like [Browse.ai](#)) with automation tools (like Make or Zapier) to push raw data straight into Notion or Airtable.

Final Thoughts on Training Your AI Dragon

If you are building out workflows, content, or systems around "training your AI dragon," keep these three fundamental operational rules in mind:

1. Stop Yelling Commands (Build the Harness)

New users treat AI like an unruly beast—they throw vague commands at it ("Write me a cover letter!") and get frustrated when it produces fire-breathing, hallucinated slop.

Training the dragon means building structured frameworks. Give it a Fact Bank (constraints), a System Prompt (identity), and a Step-by-Step Sequence (clear direction). The tighter your constraints, the more precise the execution.

2. Human in the Loop

The most iconic scene in the movie is Hiccup extending his hand, closing his eyes, and letting Toothless bridge the final gap. In AI terms, this is the Human-in-the-Loop. Never let your AI systems run entirely on autopilot. You must verify the outputs, check the data integrity of your scraped dashboards, and ensure the tailored resume still sounds like your human voice. The AI provides the speed and lift; you provide the soul and direction.

3. Change the Village Mindset

Hiccup didn't just train Toothless; he ultimately changed how the entire village operated. When you train your AI dragon to scrape data, manage personal dashboards, and systematically isolate your best career metrics, you stop acting like a desperate applicant. You start operating like a person who understands exactly how to navigate the modern digital landscape.

That's it! You've now learned everything I know about how to train your AI dragon. I hope you've enjoyed this series! You can let me know your thoughts by hitting reply or filling out this quick, 3-item [survey](#).

And if you'd like an e-book version of this series, it's free for you, as one of my first 10 subscribers. Download it [here](#).

Thanks for reading!

Sincerely,

Lisa M. Lines, PhD, MPH

LisaLines.com



Lisa Lines Consulting

Christchurch, New Zealand

You received this email because you signed up on my website.

[Unsubscribe](#) | [Update preferences](#)

